

Demonstrate the steps to build a machine-learning model that predicts the median housing price using the California housing price dataset.

1 Download the dataset : <https://media.geeksforgeeks.org/wp-content/uploads/20240319120216/housing.csv>

1. Perform the describe and info steps
2. Plot the histogram of each feature( Indicate what does histogram indicate on median\_income and house\_median\_age)
3. Demonstrate the process of creating a test set( write the difference between random and stratified test set)
4. List the geographical features from the dataset and plot a graph to Visualize Geographical Data(what does the graph indicate w.r.t housing prices and location)
5. Plot a graph to show features correlation with housing price. Which feature correlates to the maximum. Plot the graph for that with housing price and analyze what the graph indicate
6. List the features that could be combined to improve correlation and plot again to see if correlation has improved
7. List the features that needs to be cleaned and demonstrate the process of cleaning
8. Is there any categorical data that needs to be converted to numerical? If so explain the method used to convert and code the same and show the output.
9. Discuss the importance of feature scaling
10. Design a pipeline inculcating (Custom transform, feature scaling and encoding). Explain how it works

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

## 1)Perform the describe and info steps

```
import pandas as pd
```

```
housing = pd.read_csv("/content/drive/MyDrive/MLlab dataset/housing10112025.csv")
print(housing.info())
print(housing.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   longitude              20640 non-null float64
1   latitude               20640 non-null float64
2   housing_median_age     20640 non-null float64
3   total_rooms            20640 non-null float64
4   total_bedrooms         20433 non-null float64
5   population              20640 non-null float64
6   households              20640 non-null float64
7   median_income           20640 non-null float64
8   median_house_value      20640 non-null float64
9   ocean_proximity         20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
None
```

	longitude	latitude	housing_median_age	total_rooms
count	20640.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081
std	2.003532	2.135952	12.585558	2181.615252
min	-124.350000	32.540000	1.000000	2.000000
25%	-121.800000	33.930000	18.000000	1447.750000
50%	-118.490000	34.260000	29.000000	2127.000000
75%	-118.010000	37.710000	37.000000	3148.000000
max	-114.310000	41.950000	52.000000	39320.000000

	total_bedrooms	population	households	median_income
count	20433.000000	20640.000000	20640.000000	20640.000000
mean	537.870553	1425.476744	499.539680	3.870671
std	421.385070	1132.462122	382.329753	1.899822
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

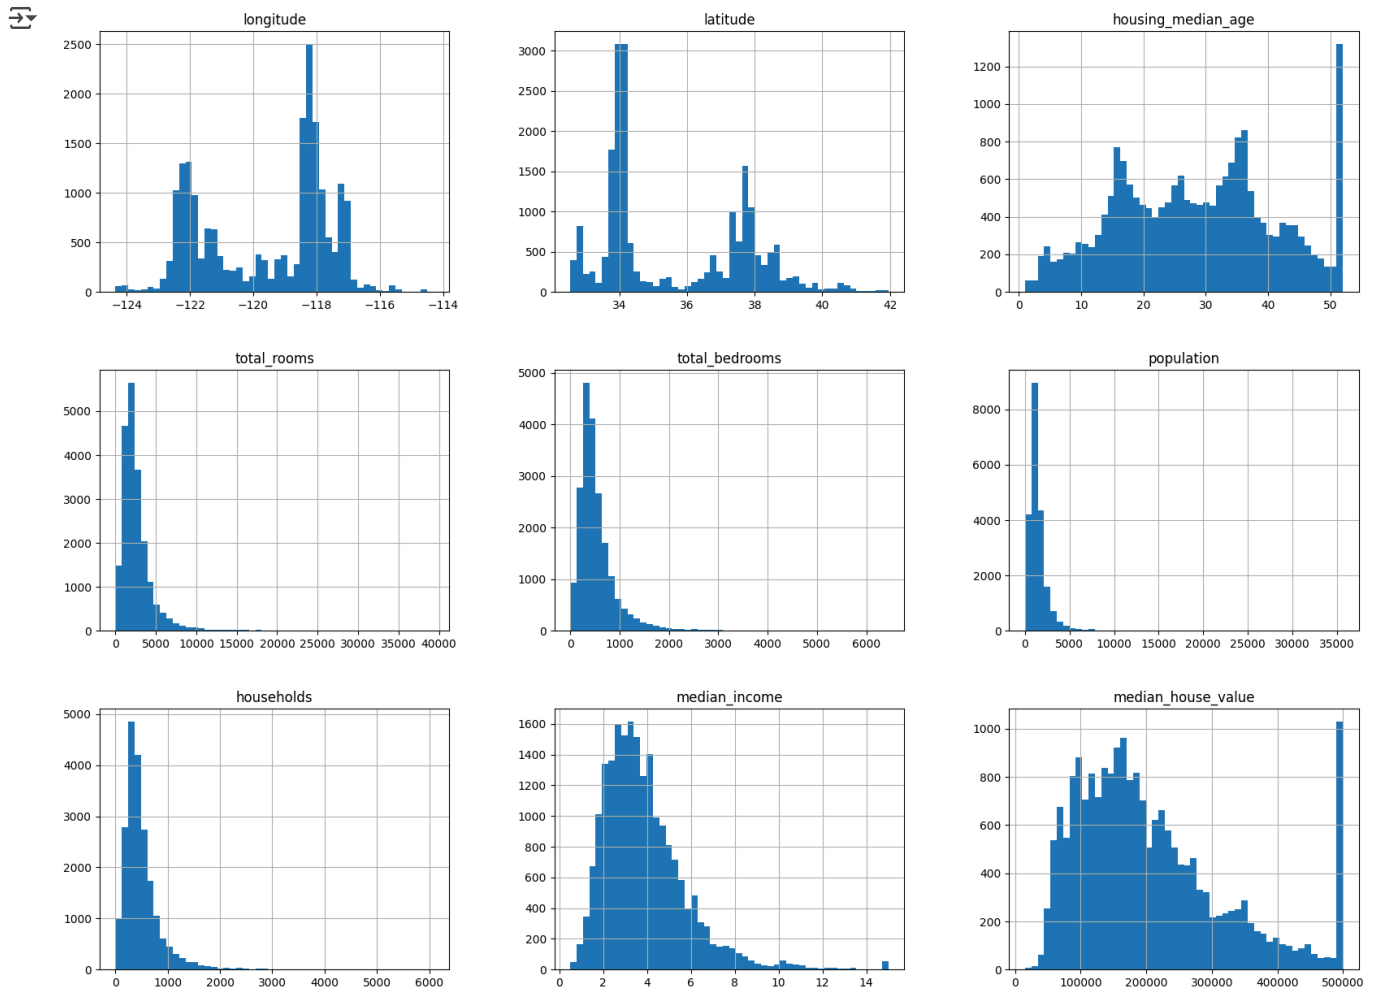
	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000

max 500001.000000

- 2) Plot the histogram of each feature( Indicate what does histogram indicate on median\_income and house\_median\_age)

```
import matplotlib.pyplot as plt
```

```
# Plot histogram for each feature
housing.hist(bins=50, figsize=(20, 15))
plt.show()
```



The histograms for median\_income and house\_median\_age can give us insights into the distribution of these features. For example, median\_income might show a right-skewed distribution, indicating that most households have lower median incomes, while

house\_median\_age might show how the ages of houses are distributed

3)Demonstrate the process of creating a test set( write the difference between random and stratified test set)

Start coding or generate with AI.

random sampling

```
from sklearn.model_selection import train_test_split

train_set, test_set = train_test_split(housing, test_size=0.2, random_state=42)
print(len(train_set))
print(len(test_set))
```

16512
4128

comparison between Random and Stratified Test Sets in tabular form:

Feature	Random Test Set	Stratified Test Set
Definition	Splits the data randomly into training and testing sets.	Ensures the training and testing sets are representative of the overall distribution of a particular feature.
Data Distribution	May not preserve the distribution of key features in training and testing sets.	Preserves the distribution of key features (e.g., income levels) across the training and testing sets.
Bias	May introduce bias due to uneven feature distribution.	Reduces bias by maintaining consistent feature distribution.
Use Case	Suitable for general datasets without specific feature distribution requirements.	Suitable for datasets with significant feature distributions that need to be maintained.

stratified approach

```
from sklearn.model_selection import train_test_split

X = housing.drop(['median_house_value'], axis=1)
y = housing['median_house_value']

train_X, test_X, train_y, test_y = train_test_split(X, y, test_size=0.2, random_state=42, stratify=housing['ocean_proximity'])
```

Double-click (or enter) to edit

proper approach

```
from sklearn.model_selection import train_test_split

# Split the data into training and testing sets using random sampling
train_set_random, test_set_random = train_test_split(housing, test_size=0.2, random_state=42)


from sklearn.model_selection import StratifiedShuffleSplit
import numpy as np
# Create a new column 'income_cat' to stratify by
housing["income_cat"] = pd.cut(housing["median_income"],
                               bins=[0., 1.5, 3.0, 4.5, 6., np.inf],
                               labels=[1, 2, 3, 4, 5])

# Use StratifiedShuffleSplit to create training and testing sets
split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
for train_index, test_index in split.split(housing, housing["income_cat"]):
    strat_train_set = housing.loc[train_index]
    strat_test_set = housing.loc[test_index]

# Drop the 'income_cat' column to return to the original state
for set_ in (strat_train_set, strat_test_set):
    set_.drop(["income_cat"], axis=1, inplace=True)
```

- ✓ 4) List the geographical features from the dataset and plot a graph to Visualize Geographical Data (what does the graph indicate w.r.t housing prices and location)

```
housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.4,
              s=housing["population"]/100, label="population", figsize=(10,7),
              c="median_house_value", cmap=plt.get_cmap("jet"), colorbar=True,
              sharex=False)
plt.legend()
plt.show()
```



This graph visualizes housing prices in relation to their geographic locations, with the color representing the median house value and the size of the circle representing the population.

- ✓ 6) Plot a graph to show features correlation with housing price. Which feature correlates to the maximum. Plot the graph for that with housing price and analyze what the graph indicates

```
# Drop the 'ocean_proximity' column
housing_num = housing.drop("ocean_proximity", axis=1)


# Calculate and print correlation matrix
corr_matrix = housing_num.corr()
print(corr_matrix["median_house_value"].sort_values(ascending=False))

# Plot the correlation between 'median_income' and 'median_house_value'
housing.plot(kind="scatter", x="median_income", y="median_house_value", alpha=0.1)
plt.show()
```

```

median_house_value    1.000000
median_income         0.688075
income_cat            0.643892
total_rooms           0.134153
housing_median_age    0.105623
households            0.065843
total_bedrooms        0.049686
population            -0.024650
longitude             -0.045967
latitude              -0.144160
Name: median_house_value, dtype: float64

```



median\_income usually shows the highest correlation with the median house value.

7) List the features that could be combined to improve correlation and plot again to see if correlation has improved

```

housing["rooms_per_household"] = housing["total_rooms"] / housing["households"]
housing["bedrooms_per_room"] = housing["total_bedrooms"] / housing["total_rooms"]

```

```

# Drop the 'ocean_proximity' column for correlation calculation
housing_num = housing.drop("ocean_proximity", axis=1)

```

```

# Calculate the new correlation matrix
corr_matrix = housing_num.corr()

```

```

# Print the sorted correlations with 'median_house_value'
print(corr_matrix["median_house_value"].sort_values(ascending=False))

```

```

median_house_value    1.000000
median_income         0.688075
income_cat            0.643892
rooms_per_household   0.151948
total_rooms           0.134153
housing_median_age    0.105623
households            0.065843
total_bedrooms        0.049686
population            -0.024650
longitude             -0.045967
latitude              -0.144160
bedrooms_per_room     -0.255880
Name: median_house_value, dtype: float64

```

```

import matplotlib.pyplot as plt

```

```

# Plot 'rooms_per_household' vs. 'median_house_value'
housing.plot(kind="scatter", x="rooms_per_household", y="median_house_value", alpha=0.1)
plt.title("Rooms per Household vs. Median House Value")
plt.show()

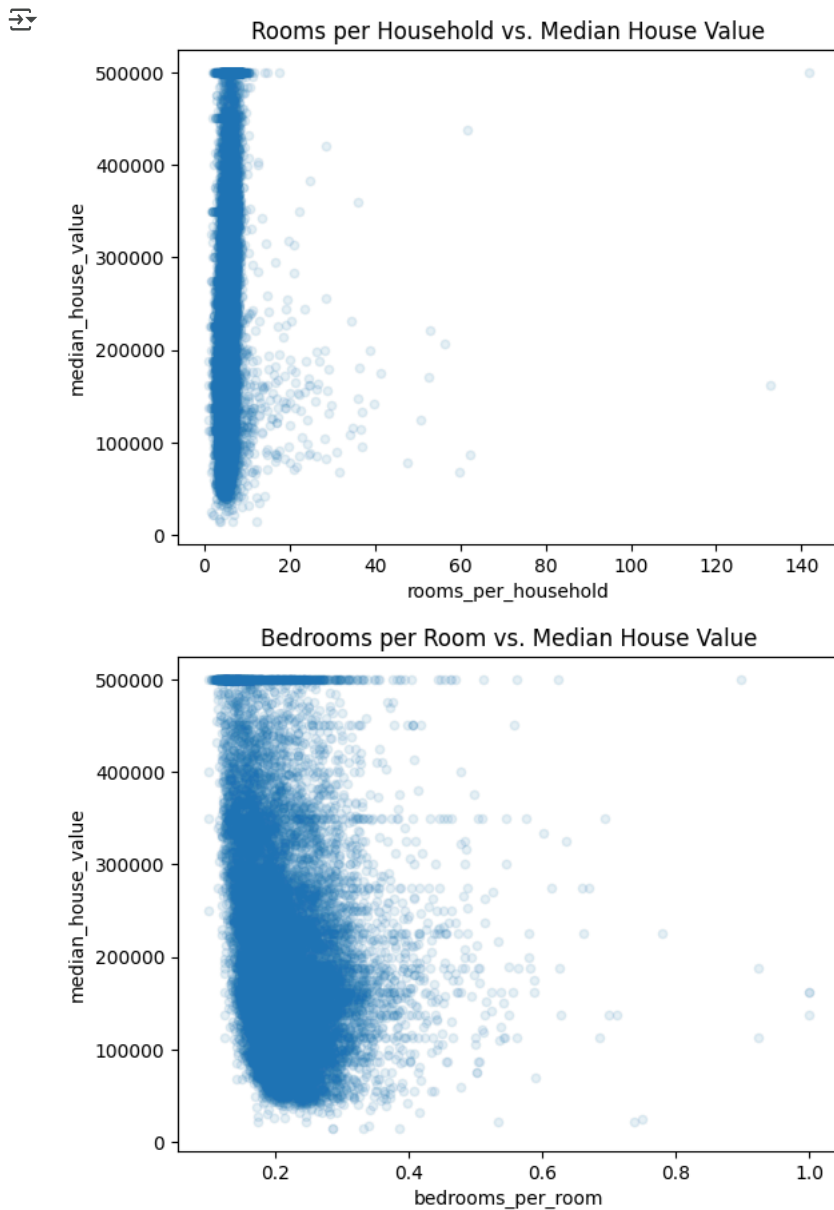
```

```

# Plot 'bedrooms_per_room' vs. 'median_house_value'

```

```
housing.plot(kind="scatter", x="bedrooms_per_room", y="median_house_value", alpha=0.1)
plt.title("Bedrooms per Room vs. Median House Value")
plt.show()
```



Analysis By creating the `rooms_per_household` and `bedrooms_per_room` features, we can observe whether the correlation with the median house value has improved. Here are some points to consider during the analysis:

**Rooms per Household:** A higher `rooms_per_household` ratio might indicate larger homes, which could be positively correlated with higher house values.

**Bedrooms per Room:** A lower `bedrooms_per_room` ratio might indicate a higher number of non-bedroom spaces, which could also be indicative of more valuable homes.

By analyzing the new correlation matrix and the scatter plots, you should be able to determine if these combined features have improved the correlation with `median_house_value`.

## ✓ 7)List the features that needs to be cleaned and demonstrate the process of cleaning

```
# Checking the column names
print(housing.columns)

# Assuming 'total_bedrooms' is present and the correct name
if "total_bedrooms" in housing.columns:
    # Option 1: Drop rows with missing values in 'total_bedrooms'
    housing.dropna(subset=["total_bedrooms"], inplace=True)

    # Option 2: Drop the 'total_bedrooms' column entirely
    housing.drop("total_bedrooms", axis=1, inplace=True)
```

```
# Option 3: Fill missing values in 'total_bedrooms' with the median value
housing["total_bedrooms"].fillna(housing["total_bedrooms"].median(), inplace=True)
else:
    print("Column 'total_bedrooms' not found in the dataset.")

Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
      'population', 'households', 'median_income', 'median_house_value',
      'ocean_proximity', 'income_cat', 'rooms_per_household',
      'bedrooms_per_room'],
      dtype='object')
Column 'total_bedrooms' not found in the dataset.
```

## ✓ Output the Corrected Data

```
print(housing.head())
```

```
longitude  latitude  housing_median_age  total_rooms  population \
0   -122.23    37.88                41.0         880.0        322.0
1   -122.22    37.86                21.0        7099.0       2401.0
2   -122.24    37.85                52.0        1467.0        496.0
3   -122.25    37.85                52.0        1274.0        558.0
4   -122.25    37.85                52.0        1627.0        565.0

households  median_income  median_house_value  ocean_proximity  income_cat \
0         126.0           8.3252           452600.0         NEAR BAY          5
1        1138.0           8.3014           358500.0         NEAR BAY          5
2         177.0           7.2574           352100.0         NEAR BAY          5
3         219.0           5.6431           341300.0         NEAR BAY          4
4         259.0           3.8462           342200.0         NEAR BAY          3

rooms_per_household  bedrooms_per_room
0           6.984127           0.146591
1           6.238137           0.155797
2           8.288136           0.129516
3           5.817352           0.184458
4           6.281853           0.172096
```

## 9. Convert Categorical Data to Numerical

```
from sklearn.preprocessing import OneHotEncoder

housing_cat = housing[["ocean_proximity"]]
encoder = OneHotEncoder()
housing_cat_1hot = encoder.fit_transform(housing_cat)
print(housing_cat_1hot.toarray())
```

```
[[0. 0. 0. 1. 0.]
 [0. 0. 0. 1. 0.]
 [0. 0. 0. 1. 0.]
 ...
 [0. 1. 0. 0. 0.]
 [0. 1. 0. 0. 0.]
 [0. 1. 0. 0. 0.]]
```

10. Importance of Feature Scaling Feature scaling is crucial to ensure that all features contribute equally to the result. Techniques include StandardScaler and MinMaxScaler.

## 11. Design a Pipeline