

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT
on

Machine Learning (23CS6PCMAL)

Submitted by

Mohammed Zeeshan Umar (1BM22CS160)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU-560019

Sep-2024 to Jan-2025

B.M.S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **Mohammed Zeeshan Umar (1BM22CS160)**, who is bonafide student of **B.M.S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Lab report has been approved as it satisfies the academic requirements in respect of an Machine Learning (23CS6PCMAL) work prescribed for the said degree.

Lab faculty Incharge Name : Amrutha Assistant Professor Department of CSE, BMSCE	Dr. Kavitha Sooda Professor & HOD Department of CSE, BMSCE
-------------------------------------------------------------------------------------------	------------------------------------------------------------------

Index

Sl. No.	Date	Experiment Title	Page No.
1	21-2-2025	Write a python program to import and export data using Pandas library functions	4
2	3-3-2025	Demonstrate various data pre-processing techniques for a given dataset	7
3	10-3-2025	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	15
4	17-3-2025	Build Logistic Regression Model for a given dataset	26
5	24-3-2025	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	30
6	7-4-2025	Build KNN Classification model for a given dataset.	35
7	21-4-2025	Build Support vector machine model for a given dataset	48
8	5-5-2025	Implement Random forest ensemble method on a given dataset.	57
9	5-5-2025	Implement Boosting ensemble method on a given dataset.	61
10	12-5-2025	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	64
11	12-5-2025	Implement Dimensionality reduction using Principal Component Analysis (PCA) method.	71

Program 1

Write a python program to import and export data using Pandas library functions

Screenshot

The image shows two pages of handwritten notes and code. The left page is titled 'Lab-1' and contains a list of tasks and a corresponding Python code snippet. The right page contains answers to three questions about handling missing values, categorical data, and feature scaling.

Lab-1

Write python code, consider filename as "housing.csv"

- To load .csv file into the data frame.
- To display information of all columns.
- To display statistical information of all numerical.
- To display the count of unique labels for "Ocean Proximity" column.
- To display which attributes (columns) in a dataset have missing value count greater than zero.

```
(i). df = pd.read_csv("housing.csv")
(ii) print(df)
(iii) print(df.describe())
(iv). print(df["Ocean Proximity"].value_count())
(v). missing_value = df.isnull().sum()
     missing_column = df[missing_value[missing_value > 0].index]
     print(missing_column)
```

1. Which column in the dataset had missing values? How did you handle them?

Dropped the missing values with the help of ~~drop na~~ `dropna` since the number of rows are less with missing value is less compatibility.

2. What were the categorical data columns did you identify in the dataset? How did you encode them?

diabetes.csv: ["Gender", "Class"]
adult.csv: ["workclass", "education", "marital status", "occupation", "relationship", "gender", "native-country", "income"]

and we used `LabelEncoder` to encode it.

3. What is the difference between Min-Max Scaling and standardization? When would you use one over the other?

Both MinMax Scaling and standardization are feature scaling techniques.

$$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}} \quad x' = \frac{x - \mu}{\sigma}$$

we use min max when we do not have significant outlier.
and we use standardization when we have outliers.

Code:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from scipy import stats

***Diabetes Dataset**
df=pd.read_csv('/content/Dataset of Diabetes .csv')
df.head()
df.shape
print(df.info())
# Summary statistics
print(df.describe())
missing_values=df.isnull().sum()
print(missing_values[missing_values > 0])
categorical_cols = df.select_dtypes(include=['object']).columns
print("Categorical columns identified:", categorical_cols)
if len(categorical_cols) > 0:
    df = pd.get_dummies(df, columns=categorical_cols, drop_first=True)
    print("\nDataFrame after one-hot encoding:")
    print(df.head())
else:
    print("\nNo categorical columns found in the dataset.")
from sklearn.preprocessing import MinMaxScaler, StandardScaler
import pandas as pd

numerical_cols = df.select_dtypes(include=['number']).columns

scaler = MinMaxScaler()
df_minmax = df.copy() # Create a copy to avoid modifying the original
df_minmax[numerical_cols] = scaler.fit_transform(df[numerical_cols])

scaler = StandardScaler()
df_standard = df.copy()
df_standard[numerical_cols] = scaler.fit_transform(df[numerical_cols])
print("\nDataFrame after Min-Max Scaling:")
print(df_minmax.head())
print("\nDataFrame after Standardization:")
print(df_standard.head())

***Adult Income Dataset**
df1=pd.read_csv('/content/adult.csv')
df1.head()
df1.shape
```

```

print(df1.info())
# Summary statistics
print(df.describe())
missing_values=df1.isnull().sum()
print(missing_values[missing_values > 0])
categorical_cols = df1.select_dtypes(include=['object']).columns
print("Categorical columns identified:", categorical_cols)
if len(categorical_cols) > 0:
    df1 = pd.get_dummies(df1, columns=categorical_cols, drop_first=True)
    print("\nDataFrame after one-hot encoding:")
    print(df.head())
else:
    print("\nNo categorical columns found in the dataset.")
from sklearn.preprocessing import MinMaxScaler, StandardScaler
import pandas as pd

numerical_cols = df1.select_dtypes(include=['number']).columns

scaler = MinMaxScaler()
df_minmax = df1.copy() # Create a copy to avoid modifying the original
df_minmax[numerical_cols] = scaler.fit_transform(df1[numerical_cols])

scaler = StandardScaler()
df_standard = df1.copy()
df_standard[numerical_cols] = scaler.fit_transform(df1[numerical_cols])
print("\nDataFrame after Min-Max Scaling:")
print(df_minmax.head())
print("\nDataFrame after Standardization:")
print(df_standard.head())

```

PROGRAM 2 Demonstrate various data pre-processing techniques for a given dataset

Screenshot

10/02/25 LAB-2

Demonstrate the steps to build a machine-learning model that predicts the median housing price using the California housing price dataset.

1. Perform the describe and info steps:

```
import pandas as pd
housing = pd.read_csv("content/drive/...")
print(housing.info())
print(housing.describe())
```

Output: column = 9
rows = 20640.

2. Plot histogram for each feature

The histogram for median-income & house-median age can give us insights into the distribution of these features. For example, median-income might show a right-skewed distribution, indicating that most households have lower median income, while house-median age might show how the ages of houses are distributed.

3. Demonstrate process of creating set & difference b/w random & stratified test.

Random	Stratified
Splits data randomly into training & testing	Both are representative of the overall distribution of a feature.
May not preserve the distribution of key feature in training & testing set.	Preserves the distribution of key features across training & testing.

4. List geographical feature.

This graph visualizes housing prices in relation to their geographic location, with the color representing median house value and size represents population.

5. Which feature correlates to maximum:

Median income usually shows the highest correlation with the median house value.

6. List the feature that could improve correlation.

By creating rooms per household & bedrooms-per room.

The median house value is improved.

7. total bedrooms list the features that leads data to be cleaned & demonstrate cleaning process.

total bedrooms.

8. Categorical data to numerical data.

Yes, ocean-proximity is categorical feature and method used to convert is OneHotEncoder().

9. Important Scaling feature:-

- ensure all feature contribute equally to result.

Technique used StandardScaler & Minmax scaler.

Code

```
import pandas as pd

import numpy as np

import seaborn as sns

import matplotlib.pyplot as plt

df=pd.read_csv('housing.csv')

df.head(2)

df.describe()

df.info()

sns.histplot(df['median_income'], kde=True, color='green')

sns.histplot(df['housing_median_age'])

from sklearn.model_selection import train_test_split

X = df.drop("median_house_value", axis=1)

y = df["median_house_value"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

X = df.drop("median_house_value", axis=1)

y = df["median_house_value"]

df["income_cat"] = pd.cut(df["median_house_value"],

bins=[0, 100000, 200000, 300000, 400000, np.inf],

labels=[1, 2, 3, 4, 5])

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42,

stratify=df["income_cat"])
```



```

train_set = X_train.copy()

train_set["median_house_value"] = y_train

train_set.plot(kind="scatter", x="longitude", y="latitude", alpha=0.4,s=train_set["population"]/100,
label="population",figsize=(10,7), c="median_house_value", cmap=plt.get_cmap("jet"),

colorbar=True)

plt.legend()

numerical_columns = df.select_dtypes(include=['float64', 'int64'])

correlation_matrix = numerical_columns.corr()

print(correlation_matrix["median_house_value"].sort_values(ascending=False))

df.plot(kind="scatter", x="median_income", y="median_house_value", alpha=0.1)

# Combine 'median_income' and 'households'

df["income_households"] = df["median_income"] * df["households"]


numerical_columns = df.select_dtypes(include=['float64', 'int64'])

correlation_matrix = numerical_columns.corr()

print(correlation_matrix["median_house_value"].sort_values(ascending=False))

df.plot(kind="scatter", x="income_households", y="median_house_value", alpha=0.1)

plt.show()

missing_values = df.isnull().sum()

print(missing_values[missing_values > 0])

h=df

h.dropna(subset=["total_bedrooms"])

from sklearn.preprocessing import OneHotEncoder

df1=pd.read_csv('housing.csv')

hc=df1[["ocean_proximity"]]

```

```

encoder=OneHotEncoder()

hc_encoded=encoder.fit_transform(hc).toarray()

hc_1hot_df = pd.DataFrame(hc_encoded, columns=encoder.get_feature_names_out(hc.columns))

hc_1hot_df.head()

```

Feature scaling is crucial in machine learning for several reasons, particularly when using algorithms that are sensitive to the scale of features. Here's a breakdown of its importance:

1. ****Improved Performance of Distance-Based Algorithms:****
2. ****Faster Convergence of Gradient Descent:****
3. ****Improved Regularization:****
4. ****Better Interpretation of Coefficients:****
5. ****Numerical Stability:****

```

from sklearn.base import BaseEstimator, TransformerMixin

from sklearn.pipeline import Pipeline

from sklearn.compose import ColumnTransformer

from sklearn.preprocessing import StandardScaler

# Custom transformer to add engineered attributes

class CombinedAttributesAdder(BaseEstimator, TransformerMixin):

    def __init__(self, add_bedrooms_per_room=True):

```

```

self.add_bedrooms_per_room = add_bedrooms_per_room

def fit(self, X, y=None):

    return self

def transform(self, X):

    # Assumes X is a NumPy array with the following columns:

    # total_rooms (index 3), total_bedrooms (index 2), population (index 4), households (index 5)

    rooms_per_household = X[:, 3] / X[:, 5]

    population_per_household = X[:, 4] / X[:, 5]

    if self.add_bedrooms_per_room:

        bedrooms_per_room = X[:, 2] / X[:, 3]

        return np.c_[X, rooms_per_household, population_per_household, bedrooms_per_room]

    else:

        return np.c_[X, rooms_per_household, population_per_household]

# Identify numerical and categorical columns

num_attribs = df1.drop("ocean_proximity", axis=1).columns # All numeric columns

cat_attribs = ["ocean_proximity"]

# Build numerical pipeline: impute missing values, add new attributes, then scale

num_pipeline = Pipeline([

    ('imputer', SimpleImputer(strategy="median")),

    ('attribs_adder', CombinedAttributesAdder()),

    ('std_scaler', StandardScaler()),

])

```

```
# Build the full pipeline combining numerical and categorical processing
```

```
full_pipeline = ColumnTransformer([  
    ("num", num_pipeline, num_attribs),  
    ("cat", OneHotEncoder(), cat_attribs),  
)
```

```
# Process the dataset using the pipeline
```

```
housing_prepared = full_pipeline.fit_transform(housing)  
print("Shape of processed data:", housing_prepared.shape)
```

Output :

Missing values in Diabetes dataset after removal:

ID	0
No_Pation	0
Gender	0
AGE	0
Urea	0
Cr	0
HbA1c	0
Chol	0
TG	0
HDL	0
LDL	0
VLDL	0
BMI	0
CLASS	0

dtype: int64

Missing values in Adult dataset after removal:

age	0
workclass	0
fnlwgt	0
education	0
educational-num	0
marital-status	0
occupation	0
relationship	0
race	0
gender	0
capital-gain	0
capital-loss	0
hours-per-week	0
native-country	0
income	0

dtype: int64

Missing values in Diabetes dataset before removal:

ID	60
No_Pation	44
Gender	56
AGE	39
Urea	58
Cr	48
HbA1c	42
Chol	47
TG	49
HDL	54
LDL	40
VLDL	46
BMI	51
CLASS	48

dtype: int64

Missing values in Adult dataset before removal:

age	2304
workclass	2324
fnlwgt	2325
education	2403
educational-num	2392
marital-status	2446
occupation	2480
relationship	2402
race	2379
gender	2463
capital-gain	2369
capital-loss	2376
hours-per-week	2367
native-country	2382
income	2324

dtype: int64

Categorical Columns in Diabetes Dataset: Index(['Gender', 'CLASS'], dtype='object')

Categorical Columns in Adult Dataset: Index(['workclass', 'education', 'marital-status', 'occupation', 'relationship', 'race', 'gender', 'native-country', 'income'], dtype='object')

Encoded Diabetes Dataset:

	ID	No_Patien	Gender	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	\
0	502.0	17975.0	0	50.0	4.7	46.0	4.9	4.2	0.9	2.4	1.4	
1	735.0	NaN	1	26.0	4.5	62.0	4.9	3.7	NaN	NaN	2.1	
2	420.0	47975.0	0	50.0	4.7	46.0	4.9	4.2	0.9	2.4	1.4	
3	680.0	87656.0	0	50.0	4.7	46.0	4.9	4.2	0.9	2.4	1.4	
4	504.0	34223.0	1	33.0	NaN	46.0	4.9	4.9	1.0	0.8	2.0	

	VLDL	BMI	CLASS
0	0.5	24.0	0
1	0.6	23.0	0
2	0.5	24.0	0
3	0.5	24.0	0
4	NaN	21.0	0

Encoded Adult Dataset:

	age	workclass	fnlwgt	education	educational-num	marital-status	\
0	25.0	4	226802.0	1	7.0		4
1	38.0	4	89814.0	11	9.0		2
2	28.0	2	336951.0	7	12.0		2
3	44.0	4	160323.0	15	10.0		2
4	18.0	0	103497.0	15	10.0		4

	occupation	relationship	race	gender	capital-gain	capital-loss	\
0	7	3	2	1	0.0	0.0	
1	5	0	4	1	0.0	0.0	
2	11	0	4	2	0.0	NaN	
3	7	0	2	1	7688.0	0.0	
4	0	3	4	0	0.0	0.0	

	hours-per-week	native-country	income
0	40.0	39	0
1	50.0	39	0
2	40.0	39	1
3	40.0	39	1
4	30.0	39	0

	Feature	MinMax_Scaled	Standardized
0	10	0.000000	-0.502219
1	20	0.010101	-0.474574
2	30	0.020202	-0.446929
3	40	0.030303	-0.419284
4	50	0.040404	-0.391639
5	1000	1.000000	2.234643

PROGRAM 3 Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

Screenshot

LAB-3

Solving the following Linear Regression Problem using Matrix approach. Find linear Regression of the data of week and product sales.

x_i (week)	y_i (Sales in thousands)
1	2
2	4
3	5
4	9

$\beta = ((X^T X)^{-1} X^T) y$

$X^T X = \begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \end{pmatrix} \times \begin{pmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \\ 1 & 4 \end{pmatrix} = \begin{bmatrix} 4 & 10 \\ 10 & 30 \end{bmatrix}$

$(X^T X)^{-1} = \begin{bmatrix} 4 & 10 \\ 10 & 30 \end{bmatrix}^{-1} = \begin{bmatrix} 1.5 & -0.5 \\ -0.5 & 0.2 \end{bmatrix}$

$(X^T X)^{-1} X^T = \begin{bmatrix} 1.5 & -0.5 \\ -0.5 & 0.2 \end{bmatrix} \times \begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \end{pmatrix} = \begin{bmatrix} 1 & 0.5 & 0 & -0.5 \\ -0.3 & -0.1 & 0.1 & 0.3 \end{bmatrix}$

Finally,

$((X^T X)^{-1} X^T) y = \begin{bmatrix} 1 & 0.5 & 0 & -0.5 \\ -0.3 & -0.1 & 0.1 & 0.3 \end{bmatrix} \times \begin{bmatrix} 2 \\ 4 \\ 5 \\ 9 \end{bmatrix}$

$= \begin{bmatrix} -0.5 \\ -0.6 - 0.4 + 0.5 + 2.7 \end{bmatrix}$

$= \begin{bmatrix} -0.5 \\ 2.2 \end{bmatrix}$

$y = -0.5 + 2.2x$

→ Normal:

$\sum x = 10$ $\sum x^2 = 30$
 $\sum y = 20$ $\sum xy = 61$

$\bar{x} = 2.5$ $\bar{y} = 5$, $\bar{x}^2 = 7.5$, $\bar{x}\bar{y} = 15.25$

$\frac{\sum y - \bar{y}}{\sum x - (\bar{x})^2}$

$= \frac{15.25 - 5 \times 2.5}{7.5 - 6.25} = 2.2$

$x_0 = \bar{y} - \bar{x}_1 \bar{x} = 5 - 2.2 \times 2.5 = -0.5$

$y = 2.2x - 0.5$

<p>1) Did you perform any data processing steps?</p> <p>Yes, handled missing values by filling them with the column mean. Also applied label encoding to categorical (like 'state') and scaled numerical features for 1000 companies.csv to normalize the data.</p> <p>2) Did you visualize the regression line for cannada-per-capita-income.csv?</p> <p>Yes, the regression line was plotted. The plot shows a strong linear relationship between year & per-capita income, meaning that as the year increases, per capita income also rises.</p> <p>3) Predicted salary for (10 years experience, 10 test score, 10 interview score)?</p> <p>The predicted salary is printed in the script and depends on the trained model's coefficient.</p> <p>4) Did you encode categorical variables for 1000 companies.csv?</p> <p>Yes, the 'state' column was encoded using labelencoder().</p>	<p>Did you scale the feature? if yes, why?</p> <p>Yes, because of R & D Spend, Administration, & Marketing Spend have different units, feature scaling (using StandardScaler()) was applied to improve model performance.</p>
-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

Code

```
# -*- coding: utf-8 -*-
```

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn import linear_model
```

```
import matplotlib.pyplot as plt
```

```
df = pd.read_csv('/content/housing_area_price.csv')
```

```
df
```

```
# Commented out IPython magic to ensure Python compatibility.
```

```
# %matplotlib inline
```

```
plt.xlabel('area')
```

```
plt.ylabel('price')
```

```
plt.scatter(df.area,df.price,color='red',marker='+')
```



```
new_df = df.drop('price',axis='columns')
```

```
new_df
```

```
price = df.price
```

```
price
```

```
# Create linear regression object
```

```
reg = linear_model.LinearRegression()
```

```
reg.fit(new_df,price)
```

```
"""(1) Predict price of a home with area = 3300 sqr ft"""
```

```
reg.predict([[3300]])
```

```
reg.coef_
```

```
reg.intercept_
```

```
""" $Y = m * X + b$  (m is coefficient and b is intercept)"""
```

```
3300*135.78767123 + 180616.43835616432
```

```
"""(1) Predict price of a home with area = 5000 sqr ft"""
```

```
reg.predict([[5000]])

# -*- coding: utf-8 -*-

import pandas as pd

import numpy as np

from sklearn import linear_model

df = pd.read_csv('/content/homeprices_Multiple_LR.csv')

df

"""Data Preprocessing: Fill NA values with median value of a column"""

df.bedrooms.median()

df.bedrooms = df.bedrooms.fillna(df.bedrooms.median())

df

reg = linear_model.LinearRegression()

reg.fit(df.drop('price',axis='columns'),df.price)

reg.coef_

reg.intercept_
```

"""Find price of home with 3000 sqr ft area, 3 bedrooms, 40 year old"""

```
reg.predict([[3000, 3, 40]])
```

$112.06244194 \times 3000 + 23388.88007794 \times 3 + -3231.71790863 \times 40 + 221323.00186540384$

```
import pandas as pd
```

```
from sklearn.linear_model import LinearRegression
```

```
# Load the dataset
```

```
df1 = pd.read_csv('/content/canada_per_capita_income.csv')
```

```
# Prepare the data
```

```
X = df1.year.values.reshape(-1, 1) # Features (year)
```

```
y = df1['per capita income (US$)'] # Target (per capita income)
```

```
# Create and train the linear regression model
```

```
model = LinearRegression()
```

```
model.fit(X, y)
```

```
# Predict per capita income for 2020
```

```
year_2020 = [[2020]]
```

```
predicted_income = model.predict(year_2020)
```

```
print(f"Predicted per capita income for Canada in 2020: {predicted_income[0]:.2f}")
```

```

import pandas as pd

from sklearn.linear_model import LinearRegression

import matplotlib.pyplot as plt

# Load the dataset (canada_per_capita_income.csv)

df1 = pd.read_csv('/content/canada_per_capita_income.csv')

# Prepare the data

X = df1.year.values.reshape(-1, 1) # Features (year)

y = df1['per capita income (US$)'] # Target (per capita income)

# Create and train the linear regression model

model = LinearRegression()

model.fit(X, y)

# Create the plot

plt.figure(figsize=(8, 6))

plt.scatter(X, y, color='blue', label='Data Points') # Now using the correct X and y

plt.plot(X, model.predict(X), color='red', label='Regression Line')

plt.xlabel('Year')

plt.ylabel('Per Capita Income (US$)')

plt.title('Per Capita Income in Canada over Time')

plt.legend()

```

```

plt.grid(True)

plt.show()

import pandas as pd

from sklearn.linear_model import LinearRegression

from sklearn.impute import SimpleImputer

# Load the dataset

df = pd.read_csv('/content/salary.csv')

# Prepare the data

X = df.iloc[:, :-1].values # Features (years of experience)

y = df.iloc[:, 1].values # Target (salary)

# Impute missing values with the mean

imputer = SimpleImputer(strategy='mean') # Create an imputer object with strategy as mean

X = imputer.fit_transform(X) # Fit and transform the imputer on feature data 'X'

# Create and train the linear regression model

model = LinearRegression()

model.fit(X, y)

# Predict salary for 12 years of experience

years_experience = [[12]]

predicted_salary = model.predict(years_experience)

```

```

print(f"Predicted salary for 12 years of experience: {predicted_salary[0]:.2f}")

import pandas as pd

from sklearn.linear_model import LinearRegression

from sklearn.impute import SimpleImputer

# Load the dataset

df = pd.read_csv('/content/hiring.csv')

# Handle missing values

# Convert 'experience' column to numeric, replacing non-numeric with NaN
df['experience'] = pd.to_numeric(df['experience'], errors='coerce')

imputer = SimpleImputer(strategy='mean')

df['experience'] = imputer.fit_transform(df[['experience']])

df['test_score(out of 10)'] = imputer.fit_transform(df[['test_score(out of 10)']])

# Prepare the data

X = df.drop('salary($)', axis='columns')

y = df['salary($)']

# Create and train the linear regression model

model = LinearRegression()

model.fit(X, y)

```

```

# Predict salaries for the given candidates

candidate1 = [[2, 9, 6]]

candidate2 = [[12, 10, 10]]


predicted_salary1 = model.predict(candidate1)

predicted_salary2 = model.predict(candidate2)


print(f"Predicted salary for candidate 1: ${predicted_salary1[0]:.2f}")

print(f"Predicted salary for candidate 2: ${predicted_salary2[0]:.2f}")

import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.preprocessing import LabelEncoder, OneHotEncoder

from sklearn.compose import ColumnTransformer


# Load the dataset

df = pd.read_csv('/content/1000_Companies.csv')


# Separate features (X) and target (y)

X = df.iloc[:, :-1].values

y = df.iloc[:, 4].values


# Encode categorical data (State)

```

```

labelencoder = LabelEncoder()

X[:, 3] = labelencoder.fit_transform(X[:, 3])

ct = ColumnTransformer(
    transformers=[('encoder', OneHotEncoder(), [3])],
    remainder='passthrough'
)

X = ct.fit_transform(X)

# Avoid dummy variable trap (remove one encoded column)
X = X[:, 1:]


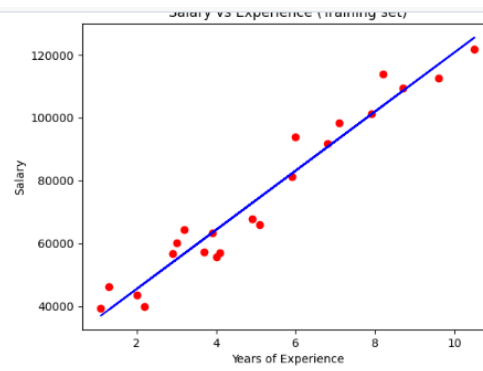

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)

# Create and train the multiple linear regression model
regressor = LinearRegression()
regressor.fit(X_train, y_train)

# Predict profit for the given values
new_prediction = regressor.predict([[1, 0, 91694.48, 515841.3, 11931.24]])
print(f"Predicted Profit: {new_prediction[0]:.2f}")

```


Output :



PROGRAM 4 Build Logistic Regression Model for a given dataset

Screenshot

24/02/25 Lab-4

Logistic Regression

(1) Student - passed or fail based on the study
Intercept $a_0 = -5$ coefficient $a_1 = 0.8$

(a) $P(z) = \frac{1}{1 + e^{-z}}$ $z = a_0 + a_1 x$
 $z = -5 + 0.8x$

(b) Probability that student who studies for 7 hours will pass
 $P(7) = \frac{e^{0.6456}}{1 + e^{0.6456}} = 0.6456$

(c) Determine predicted class (P/F) for student based on threshold 0.5
As $0.6456 > 0.5$ this student will Pass

Q2. Consider $z = [2, 1, 0]$ for 3 classes. Apply softmax function to find probability value of 3 classes.

$\text{Softmax}(z_k) = \frac{e^{z_k}}{\sum_j e^{z_j}}$
 $P(2) = \frac{e^2}{e^2 + e^1 + e^0} = 0.652$

$P(1) = \frac{e^1}{e^0 + e^1 + e^2} = 0.0900$

$P(0) = \frac{e^0}{e^0 + e^1 + e^2} = 0.2448$

For HR comma sep csv.

Q1 Which variable had clear and direct impact on employee retention? why?

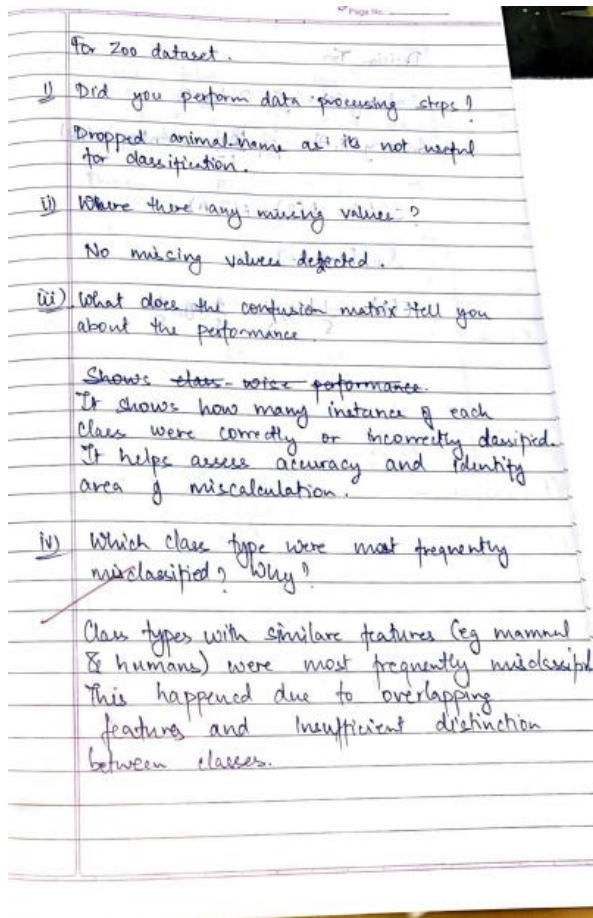
Sol: satisfaction level, time spent company, number project, average monthly hours.

Reason: They showed strong correlation with retention based on EPA.

Q2 Accuracy of logistic regression model is Good?

Accuracy (eg ~ 79%)

Yes it has good accuracy, it is above baseline, through improvement may be possible.



Code

```
import pandas as pd

import numpy as np

df=pd.read_csv("/content/HR_comma_sep.csv")

df.head(3)

print(df.isnull().sum())

print(df.groupby('left').mean(numeric_only=True))

print(df.groupby('salary').mean(numeric_only=True))

import matplotlib.pyplot as plt

pd.crosstab(df.salary,df.left).plot(kind='bar')

plt.title('Employee Retention vs Salary')
```

```

plt.xlabel('Salary')

plt.ylabel('Number of Employees')

plt.show()

pd.crosstab(df.Department,df.left).plot(kind='bar')

plt.title('Employee Retention vs Department')

plt.xlabel('Department')

plt.ylabel('Number of Employees')

plt.show()

salary_dummies = pd.get_dummies(df.salary, prefix="salary")

dept_dummies = pd.get_dummies(df.Department, prefix="dept")


df_with_dummies = pd.concat([df, salary_dummies, dept_dummies], axis=1)


df_with_dummies = df_with_dummies.drop(['salary', 'Department'], axis=1)


X_features = ['satisfaction_level', 'last_evaluation', 'number_project', 'average_monthly_hours',
'time_spend_company', 'Work_accident', 'promotion_last_5years'] + list(salary_dummies.columns) +
list(dept_dummies.columns)

X = df_with_dummies[X_features]

y = df_with_dummies.left


from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)

```

```
from sklearn.linear_model import LogisticRegression
```

```
model = LogisticRegression()
```

```
model.fit(X_train, y_train)
```

```
from sklearn.metrics import accuracy_score
```

```
y_pred = model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print("Accuracy of the model:", accuracy)
```

Output :

```

satisfaction_level  last_evaluation  number_project  average_monthly_hours  \
0                0.38              0.53              2                157
1                0.80              0.86              5                262
2                0.11              0.88              7                272
3                0.72              0.87              5                223
4                0.37              0.52              2                159

time_spend_company  Work_accident  left  promotion_last_5years  Department  \
0                 3              0      1                  0         sales
1                 6              0      1                  0         sales
2                 4              0      1                  0         sales
3                 5              0      1                  0         sales
4                 3              0      1                  0         sales



salary
0      low
1    medium
2    medium
3      low
4      low
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0   satisfaction_level    14999 non-null  float64
1   last_evaluation      14999 non-null  float64
2   number_project       14999 non-null  int64
3   average_monthly_hours 14999 non-null  int64
4   time_spend_company   14999 non-null  int64
5   Work_accident        14999 non-null  int64
6   left                 14999 non-null  int64
7   promotion_last_5years 14999 non-null  int64
8   Department           14999 non-null  object
9   salary               14999 non-null  object
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB
None

satisfaction_level  last_evaluation  number_project  \
count      14999.000000      14999.000000      14999.000000
mean         0.612834         0.716102         3.803054
std          0.240631         0.171169         1.232592
min          0.090000         0.360000         2.000000
25%          0.440000         0.560000         3.000000
50%          0.640000         0.720000         4.000000
75%          0.820000         0.870000         5.000000
max          1.000000         1.000000         7.000000

average_monthly_hours  time_spend_company  Work_accident  left  \
count      14999.000000      14999.000000      14999.000000      14999.000000
mean        201.050337         3.498233         0.144610         0.238883
std         49.943099         1.460136         0.351719         0.425924
min         96.000000         2.000000         0.000000         0.000000
25%        156.000000         3.000000         0.000000         0.000000
50%        200.000000         3.000000         0.000000         0.000000
75%        245.000000         4.000000         0.000000         0.000000
max        310.000000        10.000000         1.000000         1.000000

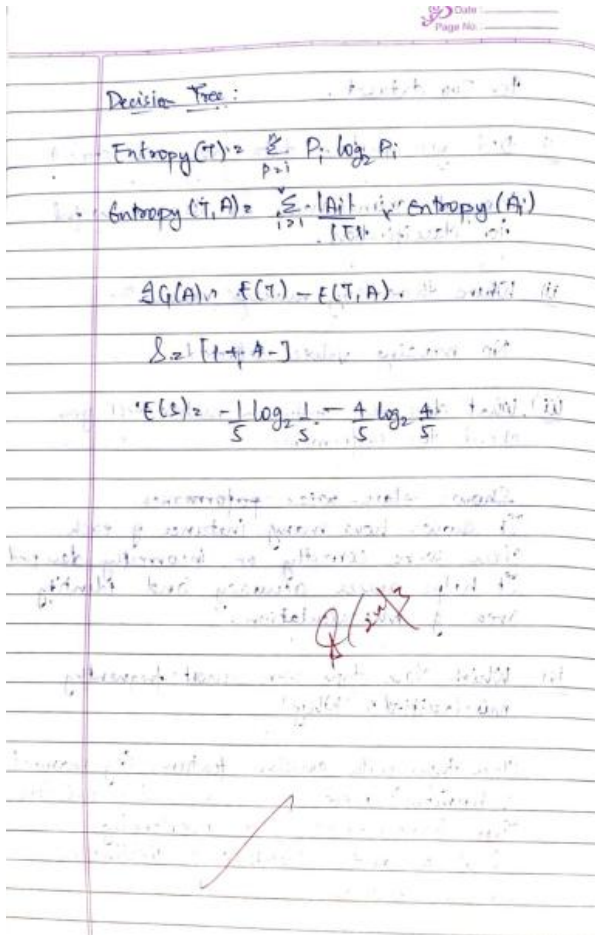
promotion last 5years

```



PROGRAM 5 Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Screenshot



Code

```
from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import accuracy_score, confusion_matrix

from sklearn import tree

import matplotlib.pyplot as plt
```

```
iris = load_iris()

X = iris.data

y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

clf = DecisionTreeClassifier()

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)

print("Accuracy:", accuracy)

print("Confusion Matrix:\n", conf_matrix)

plt.figure(figsize=(12, 8))

tree.plot_tree(clf, feature_names=iris.feature_names, class_names=iris.target_names, filled=True)

plt.show()

from sklearn.datasets import load_iris
```

```
from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import accuracy_score, confusion_matrix

from sklearn import tree

import matplotlib.pyplot as plt


iris = load_iris()

X = iris.data

y = iris.target


X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


clf = DecisionTreeClassifier()


clf.fit(X_train, y_train)


y_pred = clf.predict(X_test)


accuracy = accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)


print("Accuracy:", accuracy)

print("Confusion Matrix:\n", conf_matrix)
```



```
plt.figure(figsize=(12, 8))

tree.plot_tree(clf, feature_names=iris.feature_names, class_names=iris.target_names, filled=True)

plt.show()
```

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeRegressor

from sklearn.metrics import mean_absolute_error, mean_squared_error

import numpy as np # import numpy
```

```
data = pd.read_csv("petrol_consumption.csv")
```

```
X = data[['Petrol_tax', 'Average_income', 'Paved_Highways',
          'Population_Driver_licence(%)']]

y = data['Petrol_Consumption']
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)
```

```
regressor = DecisionTreeRegressor()
```

```
regressor.fit(X_train, y_train)
```

```

y_pred = regressor.predict(X_test)

mae = mean_absolute_error(y_test, y_pred)

mse = mean_squared_error(y_test, y_pred)

rmse = np.sqrt(mse)

print("Mean Absolute Error:", mae)

print("Mean Squared Error:", mse)

print("Root Mean Squared Error:", rmse)

from sklearn.tree import plot_tree

import matplotlib.pyplot as plt

plt.figure(figsize=(15, 10))

# Assuming 'data' is your original pandas DataFrame

plot_tree(regressor, feature_names=data[['Petrol_tax', 'Average_income', 'Paved_Highways',
'Population_Driver_licence(%)']].columns, filled=True, rounded=True)

plt.show()

```

Output :

```

Iris Dataset Results:
Accuracy: 1.00
Confusion Matrix:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
Classification Report:

```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	10
Iris-versicolor	1.00	1.00	1.00	9
Iris-virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```

Drug Dataset Results:
Accuracy: 1.00
Confusion Matrix:
[[ 6  0  0  0  0]
 [ 0  3  0  0  0]
 [ 0  0  5  0  0]
 [ 0  0  0 11  0]
 [ 0  0  0  0 15]]
Classification Report:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6
1	1.00	1.00	1.00	3
2	1.00	1.00	1.00	5
3	1.00	1.00	1.00	11
4	1.00	1.00	1.00	15
accuracy			1.00	40
macro avg	1.00	1.00	1.00	40
weighted avg	1.00	1.00	1.00	40

PROGRAM 6 Build KNN Classification model for a given dataset.

Screenshot

Lab-6
K-Nearest Neighbour

Q. Consider following dataset, $k=3$ and test data $(x, 30, 100)$ as $(Person, Age, Salary)$ solve using KNN classifier and predict target.

Person	Age	Salary	Target	Dist
A	25	50	N	52.81
B	28	55	N	46.57
C	24	40	N	31.95
D	41	60	Y	40.44
E	43	70	Y	31.04
F	38	40	Y	60.07

Test data: $(x, 30, 100)$

Nearest neighbours \rightarrow C, D, E

For class Y:

Scored: $\frac{1}{3} \times (d(Y, f(C)) + d(Y, f(D)) + d(Y, f(E)))$

Scoring: $\frac{1}{3} \times (d(N, Y) + d(N, N) + d(N, Y))$

$0 + 1 + 0 = 1$

Target-1

Q. For iris dataset, how to choose the k value? Demonstrate using accuracy rate and error rate.

Ans. Cross validation is used to find the best k value depending on accuracy or number of misclassification.

Q. For diabetes dataset, what is the purpose of feature scaling? How to perform it?

Ans. Feature scaling is used in KNN as we have the same scale because if we use a different scale, the features with larger scale get more weightage so the following lower scale features will get suppressed and won't have an impact on the model. To perform scaling, first the frequency distributions of all features were checked, the one having normal distributions were scaled using standard scaler and the rest using minmax scaler.

Code

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

import seaborn as sns

import matplotlib.pyplot as plt
```

```

try:

    data = pd.read_csv('/content/iris (1).csv')

except FileNotFoundError:

    print("Error: 'iris.csv' not found. Please upload the file to your Colab environment.")

    exit()


X = data.drop('species', axis=1)

y = data['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

knn = KNeighborsClassifier(n_neighbors=3)

knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("\nConfusion Matrix:")

cm = confusion_matrix(y_test, y_pred)

print(cm)


plt.figure(figsize=(8, 6))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',

            xticklabels=knn.classes_, yticklabels=knn.classes_)

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix')

plt.show()

```

```

print("\nClassification Report:")

print(classification_report(y_test, y_pred))


import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

from sklearn.preprocessing import StandardScaler

import seaborn as sns

import matplotlib.pyplot as plt


try:

    diabetes = pd.read_csv('diabetes.csv')

except FileNotFoundError:

    print("Error: 'diabetes.csv' not found. Please ensure the file is in the current directory.")

    exit()


X = diabetes.drop('Outcome', axis=1)

y = diabetes['Outcome']


scaler = StandardScaler()

X = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```

knn = KNeighborsClassifier(n_neighbors=5)

knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy: {accuracy}")

cm = confusion_matrix(y_test, y_pred)

print("Confusion Matrix:")

print(cm)

sns.heatmap(cm, annot=True, fmt="d")

plt.title('Confusion Matrix')

plt.xlabel('Predicted')

plt.ylabel('True')

plt.show()


print("Classification Report:")

print(classification_report(y_test, y_pred))


import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.preprocessing import StandardScaler

```

```

try:

    heart = pd.read_csv('heart.csv')

except FileNotFoundError:

    print("Error: 'heart.csv' not found. Please ensure the file is in the current directory.")

    exit()


X = heart.drop('target', axis=1)

y = heart['target']


scaler = StandardScaler()

X = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

best_k = 1

best_accuracy = 0


for k in range(1, 21):

    knn = KNeighborsClassifier(n_neighbors=k)

    knn.fit(X_train, y_train)

    y_pred = knn.predict(X_test)

    accuracy = accuracy_score(y_test, y_pred)

    if accuracy > best_accuracy:

        best_accuracy = accuracy

        best_k = k

```

```

print(f"Best k: {best_k} with accuracy {best_accuracy}")

knn = KNeighborsClassifier(n_neighbors=best_k)

knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy: {accuracy}")


cm = confusion_matrix(y_test, y_pred)

print("Confusion Matrix:")

print(cm)


sns.heatmap(cm, annot=True, fmt="d")

plt.title('Confusion Matrix')

plt.xlabel('Predicted')

plt.ylabel('True')

plt.show()


print("Classification Report:")

print(classification_report(y_test, y_pred))


import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.metrics import classification_report, confusion_matrix

```



```

cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(8, 6))

sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")

plt.title("Confusion Matrix")

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.show()

print(classification_report(y_test, y_pred))


# prompt: For Iris dataset

# How to choose the k value? Demonstrate using accuracy rate and error

# rate. Give theory


import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.preprocessing import StandardScaler


# Load the Iris dataset

try:

```

```

data = pd.read_csv('/content/iris (1).csv')

except FileNotFoundError:

    print("Error: 'iris (1).csv' not found. Please upload the file to your Colab environment.")

    exit()


# Prepare the data

X = data.drop('species', axis=1)

y = data['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Scale the data (important for KNN)

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)


# Find the optimal k value

error_rates = []

for k in range(1, 31): # Test k values from 1 to 30

    knn = KNeighborsClassifier(n_neighbors=k)

    knn.fit(X_train, y_train)

    y_pred = knn.predict(X_test)

    error_rates.append(1 - accuracy_score(y_test, y_pred)) # Error rate = 1 - accuracy


# Plot error rates

```

```

plt.figure(figsize=(10, 6))

plt.plot(range(1, 31), error_rates, color='blue', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)

plt.title('Error Rate vs. K Value')

plt.xlabel('K')

plt.ylabel('Error Rate')

plt.show()

# Theory for choosing k:

# The optimal 'k' value minimizes the error rate.

# Very small k (e.g., 1) can lead to overfitting, being too sensitive to noise.

# Very large k (e.g., 30) can lead to underfitting, smoothing out the decision boundaries too much.

# We seek a k that balances these extremes, as shown by the error rate plot.

#Select k based on the minimum error rate observed in the plot

best_k = error_rates.index(min(error_rates)) + 1 #Add 1 as the index starts from 0

# Train and evaluate the model with the best k

knn = KNeighborsClassifier(n_neighbors=best_k)

knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

```

```

# Evaluate the model

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("\nConfusion Matrix:")

cm = confusion_matrix(y_test, y_pred)

print(cm)

print("\nClassification Report:")

print(classification_report(y_test, y_pred))


plt.figure(figsize=(8, 6))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=knn.classes_, yticklabels=knn.classes_)

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix')

plt.show()


import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import accuracy_score

import matplotlib.pyplot as plt


# Load data

df = pd.read_csv('/content/iris (1).csv')

```

```

X = df.iloc[:, :-1]

y = df.iloc[:, -1]


# Train-test split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)


# Store accuracy and error rate

accuracy = []

error_rate = []


# Try k from 1 to 20

for k in range(1, 21):

    knn = KNeighborsClassifier(n_neighbors=k)

    knn.fit(X_train, y_train)

    preds = knn.predict(X_test)

    acc = accuracy_score(y_test, preds)

    accuracy.append(acc)

    error_rate.append(1 - acc)


# Plot

plt.figure(figsize=(10,5))

plt.plot(range(1, 21), accuracy, label='Accuracy')

plt.plot(range(1, 21), error_rate, label='Error Rate')

plt.xlabel('K Value')

```

```

plt.ylabel('Rate')

plt.title('K vs Accuracy and Error Rate')

plt.legend()

plt.show()


import pandas as pd

from sklearn.preprocessing import StandardScaler


# Load data

df = pd.read_csv('/content/diabetes.csv')

X = df.drop('Outcome', axis=1) # Features

y = df['Outcome']             # Target


# Perform scaling

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)


# Convert back to DataFrame (optional)

X_scaled_df = pd.DataFrame(X_scaled, columns=X.columns)

Output :

```

IRIS Dataset - Actual vs Predicted:

	Actual	Predicted
0	versicolor	versicolor
1	setosa	setosa
2	virginica	virginica
3	versicolor	versicolor
4	versicolor	versicolor
5	setosa	setosa
6	versicolor	versicolor
7	virginica	virginica
8	versicolor	versicolor
9	versicolor	versicolor
10	virginica	virginica
11	setosa	setosa
12	setosa	setosa
13	setosa	setosa
14	setosa	setosa
15	versicolor	versicolor
16	virginica	virginica
17	versicolor	versicolor
18	versicolor	versicolor
19	virginica	virginica
20	setosa	setosa
21	virginica	virginica
22	setosa	setosa
23	virginica	virginica
24	virginica	virginica
25	virginica	virginica
26	virginica	virginica
27	virginica	virginica
28	setosa	setosa
29	setosa	setosa

Accuracy Score: 1.0

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

DIABETES Dataset - Actual vs Predicted:

	Actual	Predicted
0	0	0
1	0	0
2	0	0
3	0	0
4	0	1
..
149	1	1
150	0	0
151	0	0
152	1	0
153	0	0



[154 rows x 2 columns]

Accuracy Score: 0.6948051948051948

Confusion Matrix:

```
[[79 20]
 [27 28]]
```

Best K value: 7 with Accuracy: 0.9180

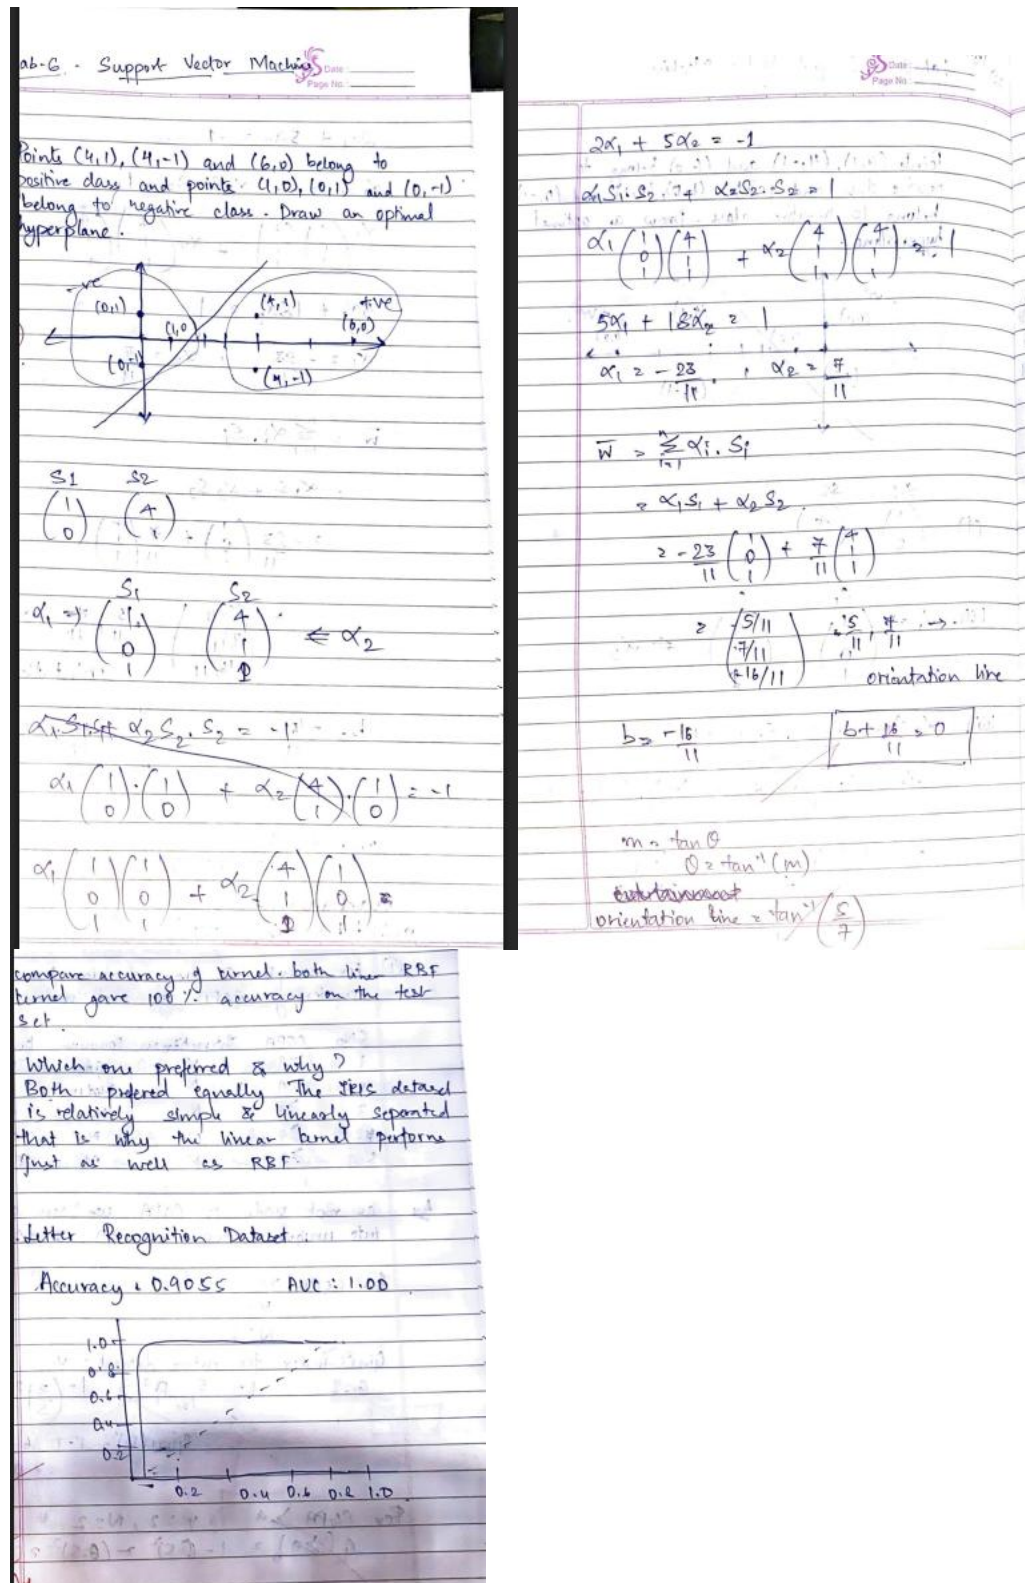


Classification Report:

	precision	recall	f1-score	support
0	0.90	0.93	0.92	29
1	0.94	0.91	0.92	32
accuracy			0.92	61
macro avg	0.92	0.92	0.92	61
weighted avg	0.92	0.92	0.92	61

PROGRAM 7 Build Support vector machine model for a given dataset

Screenshot



Code

```
import numpy as np

import matplotlib.pyplot as plt

positive_class = np.array([[4, 1], [4, -1], [6, 0]])

negative_class = np.array([[1, 0], [0, 1], [0, -1]])

plt.figure(figsize=(8, 6))

plt.scatter(positive_class[:, 0], positive_class[:, 1], color='red', label='Positive Class', s=100,
            edgecolors='black')

plt.scatter(negative_class[:, 0], negative_class[:, 1], color='blue', label='Negative Class', s=100,
            edgecolors='black')

all_points = np.concatenate([positive_class, negative_class])

labels = ["(4,1)", "(4,-1)", "(6,0)", "(1,0)", "(0,1)", "(0,-1)"]

for i, txt in enumerate(labels):

    plt.annotate(txt, (all_points[i][0], all_points[i][1]), textcoords="offset points", xytext=(0,5),
                ha='center', fontsize=10)

x_values = np.linspace(-1, 7, 100)

y_values = np.zeros_like(x_values)

plt.plot(x_values, y_values, color='black', linestyle='--', label='Optimal Hyperplane (y = 0)')
```

```
plt.plot(x_values, y_values + 1, color='gray', linestyle=':', label='Margin at y = 1')
```

```
plt.plot(x_values, y_values - 1, color='gray', linestyle=':', label='Margin at y = -1')
```

```
plt.title('Optimal Hyperplane for SVM (Visual Approximation)', fontsize=14)
```

```
plt.xlabel('x1')
```

```
plt.ylabel('x2')
```

```
plt.xlim(-1, 7)
```

```
plt.ylim(-2, 2)
```

```
plt.axhline(0, color='black',linewidth=0.5)
```

```
plt.axvline(0, color='black',linewidth=0.5)
```

```
plt.legend()
```

```
plt.grid(True)
```

```
plt.show()
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.svm import SVC
```

```
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
data = pd.read_csv('/content/iris (1) (1).csv')
```

```

X = data.drop('species', axis=1)

y = data['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


svm_rbf = SVC(kernel='rbf')

svm_rbf.fit(X_train, y_train)

y_pred_rbf = svm_rbf.predict(X_test)

accuracy_rbf = accuracy_score(y_test, y_pred_rbf)

cm_rbf = confusion_matrix(y_test, y_pred_rbf)


print("SVM with RBF Kernel:")

print("Accuracy:", accuracy_rbf)

print("Confusion Matrix:\n", cm_rbf)


plt.figure(figsize=(6, 4))

sns.heatmap(cm_rbf, annot=True, fmt='d', cmap='Blues',

            xticklabels=data['species'].unique(),

            yticklabels=data['species'].unique())

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix (RBF Kernel)')

plt.show()


svm_linear = SVC(kernel='linear')

```

```

svm_linear.fit(X_train, y_train)

y_pred_linear = svm_linear.predict(X_test)

accuracy_linear = accuracy_score(y_test, y_pred_linear)

cm_linear = confusion_matrix(y_test, y_pred_linear)


print("\nSVM with Linear Kernel:")

print("Accuracy:", accuracy_linear)

print("Confusion Matrix:\n", cm_linear)


plt.figure(figsize=(6, 4))

sns.heatmap(cm_linear, annot=True, fmt='d', cmap='Blues',

            xticklabels=data['species'].unique(),

            yticklabels=data['species'].unique())

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix (Linear Kernel)')

plt.show()

import numpy as np

import matplotlib.pyplot as plt

import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, auc

import seaborn as sns

```

```

from sklearn.preprocessing import label_binarize

from sklearn.multiclass import OneVsRestClassifier


data = pd.read_csv('/content/letter-recognition.csv') # Replace with the correct path if necessary


X = data.drop('letter', axis=1)

y = data['letter']


X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


svm_classifier = SVC(kernel='rbf', probability=True) # probability=True is needed for ROC curve
svm_classifier.fit(X_train, y_train)


y_pred = svm_classifier.predict(X_test)


accuracy = accuracy_score(y_test, y_pred)

cm = confusion_matrix(y_test, y_pred)


print("SVM Classifier:")

print("Accuracy:", accuracy)

print("Confusion Matrix:\n", cm)


plt.figure(figsize=(10, 8))

sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=np.unique(y),
yticklabels=np.unique(y))

```

```

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix')

plt.show()


y_test_bin = label_binarize(y_test, classes=np.unique(y))

n_classes = y_test_bin.shape[1]


classifier = OneVsRestClassifier(SVC(kernel='rbf', probability=True))

classifier.fit(X_train, y_train)

y_score = classifier.predict_proba(X_test)


fpr = dict()

tpr = dict()

roc_auc = dict()

for i in range(n_classes):

    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_score[:, i])

    roc_auc[i] = auc(fpr[i], tpr[i])


fpr["micro"], tpr["micro"], _ = roc_curve(y_test_bin.ravel(), y_score.ravel())

roc_auc["micro"] = auc(fpr["micro"], tpr["micro"])

plt.figure(figsize=(8, 6))

plt.plot(fpr["micro"], tpr["micro"],

        label='micro-average ROC curve (area = {0:0.2f})'
```

```

        ".format(roc_auc["micro"]))

plt.plot([0, 1], [0, 1], 'k--')

plt.xlim([0.0, 1.0])

plt.ylim([0.0, 1.05])

plt.xlabel('False Positive Rate')

plt.ylabel('True Positive Rate')

plt.title('Micro-averaged ROC Curve')

plt.legend(loc="lower right")

plt.show()

print(f'Micro-averaged AUC: {roc_auc["micro"]}')

```

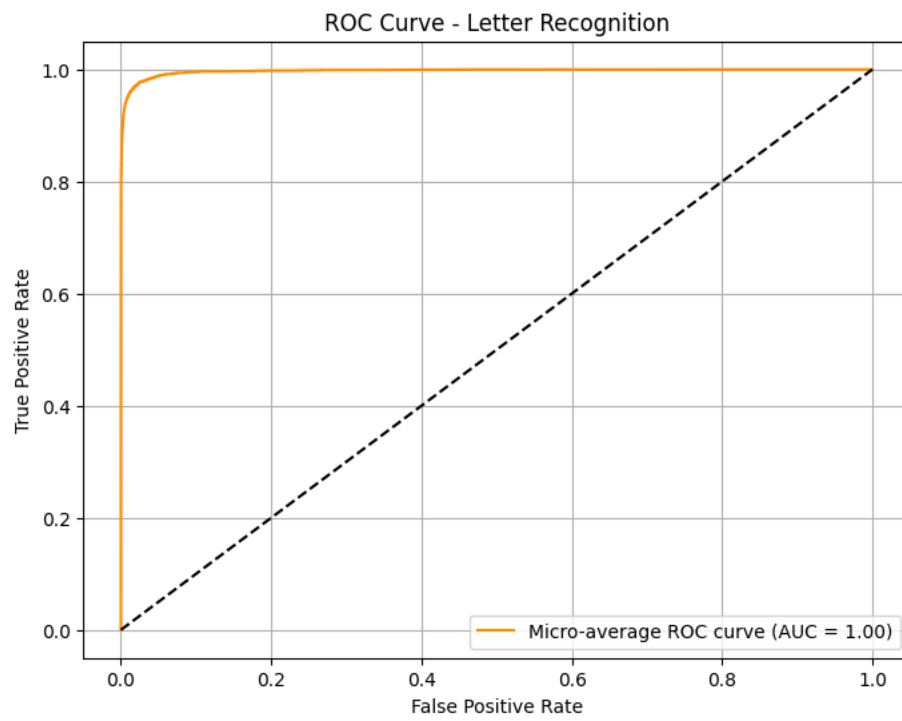
Output :

```

----- Letter Recognition Dataset -----
Accuracy: 0.95025
Confusion Matrix:
[[148  0  0  0  0  0  0  0  0  0  0  0  0  1  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0]
 [  1 147  0  3  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  2
  0  0  0  0  0  0  0  0  0]
 [  0  0 126  0  2  0  2  1  0  0  0  2  0  0  0  0  4  0  0  0
  0  0  0  0  0  0  0  0  0]
 [  0  1  0 153  0  0  0  1  0  0  0  0  0  0  1  0  0  0  0  0
  0  0  0  0  0  0  0  0  0]
 [  0  0  1  0 135  0  3  0  0  0  0  0  0  0  0  0  0  0  1  0
  0  0  0  0  0  0  0  0  1]
 [  0  1  0  0  1 134  0  0  1  0  0  0  0  0  0  0  0  1  0  0
  0  2  0  0  0  0  0  0  0]
 [  0  0  1  3  0  0 153  0  0  0  0  0  0  0  0  0  0  0  0  2
  0  0  0  1  0  0  0  0]
 [  0  3  0  4  0  0  1 118  0  0  4  0  0  1  2  0  0  0  0  8
  0  0  1  0  0  1  1  0]
 [  0  0  0  0  0  2  0  0 136  7  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  1  0  0]
 [  1  0  0  0  1  1  0  0  2 143  0  0  0  0  1  0  0  0  0  0
  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  0  0  0  1  0  0 117  0  0  0  0  0  0  0 10
  0  0  0  0  0  2  0  0]
 [  0  0  1  0  4  0  3  0  0  0  0 144  0  0  0  0  0  0  0  1
  1  0  0  0  0  0  0  1]
 [  0  2  0  0  0  0  0  0  0  0  0  0 166  0  0  0  0  0  0
  0  0  0  0  0  0  0  0]
 [  0  0  0  2  0  0  0  1  0  0  0  0  0 140  4  0  0  0  3
  0  0  0  0  1  0  0  0]
 [  0  0  0  4  0  0  1  0  0  0  0  0  0  0 136  0  1  0  0
  0  0  1  0  2  0  0  0]
 [  0  1  0  0  2 11  3  0  0  0  0  1  0  0  1 154  0  0  0
  0  0  0  0  0  0  0  0]
 [  0  1  0  1  1  0  0  0  0  0  0  0  0  0  0  0 163  0  0
  0  0  0  0  0  0  0  0]
 [  0  5  0  2  0  0  0  0  0  0  0  2  0  0  1  0  0  0 150
  0  0  0  0  0  0  0  0]
 [  0  2  0  0  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 167 0  0  0  0  0  0  1]
 [  0  1  0  0  0  0  1  0  0  0  1  0  0  0  0  0  0  0  1

```

```
[ 0 0 0 0 1 0 0 0 0 2 0 0 0 0 0 0 1  
 1 0 0 0 0 0 0 127]]
```



PROGRAM 8 Implement Random forest ensemble method on a given dataset.

Screenshot

oclosp's Lab. 7 .. Date _____
Page No. _____

Random Forest Ensemble

① For sample ST draw decision tree considering CGPA as root node

SNo	CGPA	Interactions	Common	Answered	For
1	≥ 9	M	Good	Good	Y
2	≤ 9	N	Moderate	Good	Y
3	≥ 9	N	M	Ang	N
4	≥ 9	N	M	Ang	N
5	≥ 9	Y	M	Good	Y

As root node is CGPA we can separate into unique groups as ≥ 9 & < 9 .

For $CGPA \geq 9$ Y: 2 N: 2

Gini index for entire dataset: $Y=3, N=2$

$$G.I = 1 - \sum_{i=1}^2 p_i^2 = 1 - \left(\frac{3}{5}\right)^2 - \left(\frac{2}{5}\right)^2$$

$$Gini(D) = 1 - 0.36 - 0.16 = 0.48$$

For $CGPA \geq 9$ Y: 2, N: 2

$$G(\geq 9) = 1 - (0.5)^2 - (0.5)^2 = 0.5$$

for $CGPA < 9 \rightarrow$ 1 sample
Y = 1, N = 0
 $Gini(< 9) = 0$

Weighted Gini_{CGPA} = $\frac{4}{5} \times 0.5 + \frac{1}{5} \times 0 = 0.4$

Information Gain (Gini results)
= $0.48 - 0.4 = 0.08$

Split $CGPA \geq 9$ further to get
Split on interaction Y = Yes - 2
N = N - 2

$$Gini = \frac{2}{4} \times 0 + \frac{2}{4} \times 0 = 0$$

(Gini reduction) $G = 0.5 - 0 = 0.5$

```

graph TD
    CGPA((CGPA)) -- "≥ 9" --> Interact{Interact}
    CGPA -- "< 9" --> Yes1[Yes]
    Interact -- "Yes" --> Yes2[Yes]
    Interact -- "No" --> Box1[ ]
  
```

1. $Gini = 1 - \left(\frac{4}{5}\right)^2 - \left(\frac{1}{5}\right)^2 = 1 - 0.64 - 0.04 = 0.32$

2. $Gini = \frac{2}{5} \times 0 + \frac{3}{5} \times 0.444 = 0.2664$

$\Delta Gini = 0.32 - 0.2664 = 0.0536$

3. growing the "no" branch with other features.

→ (GPA):
weighted gini = $\frac{1}{3} \times 0 + \frac{2}{3} \times 0.50 = 0.33$

$\Delta Gini = 0.11$

→ Practical Knowledge:
Weighted gini = 0
 $\Delta Gini = 0.444$ (best).

[Interactionism]

Yes: "Yes"
No: [Practical Knowledge]
 good: Yes
 Average: No

RF	n-estimators	mean accuracy
1	10	0.9667
2	50	0.9667
3	100	0.9667
4	150	0.9667
5	200	0.9667

Accuracy = 1.0
Best number of trees = 20

confusion matrix = $\begin{bmatrix} 10 & 0 & 0 \\ 0 & 9 & 0 \\ 0 & 0 & 11 \end{bmatrix}$

Code

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score

import matplotlib.pyplot as plt

# Load the dataset

df = pd.read_csv('/content/iris (1).csv')

# Prepare features and target

X = df.drop(columns=['species']) # Assuming 'species' is the target column
```

```

y = df['species']

# Split into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Build Random Forest with default n_estimators (10)

rf_default = RandomForestClassifier(n_estimators=10, random_state=42)

rf_default.fit(X_train, y_train)

y_pred_default = rf_default.predict(X_test)

# Measure accuracy

default_score = accuracy_score(y_test, y_pred_default)

print(f"Default RF accuracy (n_estimators=10): {default_score:.4f}")

# Fine-tune the number of trees

scores = []

n_range = range(1, 101)

for n in n_range:

    rf = RandomForestClassifier(n_estimators=n, random_state=42)

    rf.fit(X_train, y_train)

    y_pred = rf.predict(X_test)

    score = accuracy_score(y_test, y_pred)

    scores.append(score)

```

```

# Find the best score and number of trees

best_score = max(scores)

best_n = n_range[scores.index(best_score)]

print(f"Best RF accuracy: {best_score:.4f} with n_estimators={best_n}")

# Optional: Plot accuracy vs number of estimators

plt.figure(figsize=(10, 6))

plt.plot(n_range, scores, marker='o')

plt.title('Random Forest Accuracy vs Number of Trees')

plt.xlabel('Number of Trees (n_estimators)')

plt.ylabel('Accuracy')

plt.grid(True)

plt.show()

```

Output :

```

Default n_estimators=10 → Mean CV accuracy: 0.9667

Random Forest Tuning Results   n_estimators  mean_accuracy
0           10         0.966667
1           50         0.966667
2          100         0.966667
3          150         0.966667
4          200         0.966667
Best number of trees: 10 → Mean CV accuracy: 0.9667

```

PROGRAM 9 Implement Boosting ensemble method on a given dataset.

Screenshot

05/05/25 Lab - 8

Boosting ensemble method

CRPD	Interac	Practi	Comm	Job
<9	Y	G	G	Y
<9	N	G	M	Y
<9	N	A	M	N
<9	N	A	G	N
>9	Y	G	M	Y
>9	Y	G	M	Y

Initial weight = 1/6

$$\sum_{j=1}^5 f_j(d_j) \cdot \text{wt}(d_j)$$

$$= 2 \times \frac{1}{6} = 0.33$$

$$\alpha = \frac{1}{2} \ln \left(\frac{1 - \epsilon_{\text{weak}}}{\epsilon_{\text{weak}}} \right)$$

$$\alpha_{\text{weak}} = \frac{1}{2} \ln \left(\frac{1 - 0.33}{0.33} \right)$$

$$\alpha_{\text{weak}} = 0.347$$

$$= \frac{1}{6} \times 4 \times e^{-0.0244} \rightarrow \frac{1}{6} \times 2 \times e^{0.0244}$$

$$= 0.9428$$

$$\text{New} = 0.1249$$

$$0.2501$$

Adaboost weak classifier

Iteration	Weak Classifier	Weight	Accuracy
1	1	10	0.8203
2	2	50	0.8300
3	3	100	0.8305
4	4	150	0.8322
5	5	200	0.8326

Confusion Matrix

	Actual Y	Actual N
Predicted Y	7028	1386
Predicted N	1243	1112

Code

```
import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder

from sklearn.ensemble import AdaBoostClassifier

from sklearn.metrics import accuracy_score

from sklearn.tree import DecisionTreeClassifier
```

```

# Load dataset

df = pd.read_csv("/content/income.csv")

# Drop rows with missing values

df.dropna(inplace=True)

# Encode categorical columns

label_encoders = {}

for column in df.select_dtypes(include=['object']).columns:

    le = LabelEncoder()

    df[column] = le.fit_transform(df[column])

    label_encoders[column] = le

# Separate features and target

X = df.drop(columns=['income_level'], errors='ignore', axis=1)

y = df['income_level']

# Split into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# AdaBoost with 10 estimators

model_10 = AdaBoostClassifier(n_estimators=10, random_state=42)

model_10.fit(X_train, y_train)

y_pred_10 = model_10.predict(X_test)

score_10 = accuracy_score(y_test, y_pred_10)

print(f"Accuracy with 10 estimators: {score_10:.4f}")

# Fine-tune number of estimators

best_score = 0

best_n = 0

```

```

estimators_range = list(range(10, 201, 10))

scores = []

for n in estimators_range:

    model = AdaBoostClassifier(n_estimators=n, random_state=42)

    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)

    score = accuracy_score(y_test, y_pred)

    scores.append(score)

    print(f"n_estimators={n}, Accuracy={score:.4f}")

    if score > best_score:

        best_score = score

        best_n = n

print(f"\nBest Accuracy: {best_score:.4f} using {best_n} estimators")

# Plot accuracy vs number of estimators

plt.figure(figsize=(7, 4))

plt.plot(estimators_range, scores, marker='o', linestyle='-', color='blue')

plt.title("Accuracy vs Number of Estimators (AdaBoost)")

plt.xlabel("Number of Estimators (Trees)")

plt.ylabel("Accuracy")

plt.grid(True)

plt.xticks(estimators_range)

plt.tight_layout()

plt.show()

```

Output :

Default AdaBoost (n_estimators=10) ⇒ Mean 5-fold CV accuracy: 0.8202

Tuning results:

n_estimators	mean_accuracy
0	10 0.820237
1	50 0.830023
2	100 0.832050
3	150 0.832296
4	200 0.832624

Best performance: 0.8326 accuracy using 200 trees

PROGRAM 10 Build k-Means algorithm to cluster a set of data stored in a .CSV file.

Screenshot

10/06/25 Lab - 9

Q. Compute two clusters using k-means algorithm for clustering whose cluster centers are (1,1) & (5,7) execute for 2 Iteration.

→ Iteration 1:-

Calc. Euclidean dist to centroids:

Record	Close to C ₁	Close to C ₂	Assignment
(1,1)	0	7.21	C ₁
(1.5,2)	1.12	6.10	C ₁
(3,4)	3.61	3.61	C ₁
(5,7)	7.21	0.0	C ₂
(3.5,5)	4.12	2.5	C ₂
(4.5,5)	5.31	2.06	C ₂
(3.5,4.5)	4.30	2.92	C ₂

New centroids:

$$C_1 = \frac{1+1.5+3}{3}, \frac{1+2+4}{3} = 1.83, 2.33$$

$$C_2 = \frac{5+3.5+4.5+3.5}{4}, \frac{7+5+5+4.5}{4} = 4.12, 5.37$$

→ Iteration 2:

Record	Close to C ₁	Close to C ₂	Assignment
(1,1)	1.57	5.37	C ₁
(1.5,2)	0.97	4.27	C ₁
(3,4)	2.04	1.77	C ₂
(5,7)	5.67	1.85	C ₂
(3.5,5)	3.15	0.72	C ₂
(4.5,5)	3.78	0.53	C ₂
(3.5,4.5)	2.74	1.07	C ₂

∴ New cluster are
 $C_1 = \{R_1, R_2\}$, $C_2 = \{R_3, R_4, R_5, R_6\}$

New Centroids: $C_1 = \frac{1+1.5}{2}, \frac{1+2}{2} = 1.25, 1.5$
 $C_2 = \frac{3+5+3.5+4.5}{4}, \frac{4+7+5+4.5}{4} = 3.125, 5.25$

For Iris dataset

The elbow plot (number of clusters) shows a sharp elbow at k=3, indicating that the cluster is the normal choice for petal length and width of the Iris.

Code

```
import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler
```



```
from sklearn.cluster import KMeans

from sklearn.metrics import accuracy_score

from scipy.stats import mode

import matplotlib.pyplot as plt


# Step 1: Generate sample data and save to CSV

np.random.seed(42)

names = [f"Person_{i}" for i in range(50)]

ages = np.random.randint(20, 60, 50)

income = np.random.randint(30000, 120000, 50)


df = pd.DataFrame({'Name': names, 'Age': ages, 'Income': income})

df.to_csv("income.csv", index=False)


# Step 2: Load the data

data = pd.read_csv("income.csv")


# Drop 'Name' and extract features

X = data[['Age', 'Income']]


# Step 3: Split the data

X_train, X_test = train_test_split(X, test_size=0.2, random_state=42)


# Step 4: Perform scaling
```

```

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)


# Step 5: Plot SSE vs number of clusters (Elbow method)

sse = []

k_range = range(1, 11)

for k in k_range:

    kmeans = KMeans(n_clusters=k, random_state=42)

    kmeans.fit(X_train_scaled)

    sse.append(kmeans.inertia_)


plt.figure(figsize=(8, 4))

plt.plot(k_range, sse, marker='o')

plt.xlabel('Number of clusters')

plt.ylabel('SSE (Inertia)')

plt.title('Elbow Method For Optimal k')

plt.grid(True)

plt.show()


# Step 6: Choose optimal number of clusters (say 3) and fit model

optimal_k = 3

kmeans = KMeans(n_clusters=optimal_k, random_state=42)

kmeans.fit(X_train_scaled)

```

```

# Predict on test data

predictions = kmeans.predict(X_test_scaled)

# Note: There's no ground truth labels, but for demonstration,
# we can try assigning true clusters (via KMeans on full data)
# and see if predicted clusters align

# Fit on full data to assign pseudo-labels
full_kmeans = KMeans(n_clusters=optimal_k, random_state=42)
true_clusters = full_kmeans.fit_predict scaler.fit_transform(X))

# Align predicted clusters using majority voting (only for demonstration)
# Match predicted labels to closest true labels
def map_clusters(true_labels, pred_labels):
    labels = np.zeros_like(pred_labels)
    for i in range(optimal_k):
        mask = (pred_labels == i)
        if np.sum(mask) == 0:
            continue
        labels[mask] = mode(true_labels[mask])[0]
    return labels

mapped_preds = map_clusters(true_clusters[X_test.index], predictions)

```

```
accuracy = accuracy_score(true_clusters[X_test.index], mapped_preds)

print(f"Approximate Clustering Accuracy: {accuracy:.2f}")
```

```
import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

from sklearn.datasets import load_iris

from sklearn.cluster import KMeans

from sklearn.preprocessing import StandardScaler

from sklearn.metrics import silhouette_score
```

```
# Step 1: Load Iris dataset

iris = load_iris()

df = pd.DataFrame(iris.data, columns=iris.feature_names)

df['target'] = iris.target
```

```
# Keep only petal length and petal width

X = df[['petal length (cm)', 'petal width (cm)']].values
```

```
# Step 2: Check impact of scaling

# Try without scaling

sse_unscaled = []

for k in range(1, 11):

    kmeans = KMeans(n_clusters=k, random_state=42)
```

```

kmeans.fit(X)

sse_unscaled.append(kmeans.inertia_)


# Now scale the features

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)


sse_scaled = []

for k in range(1, 11):

    kmeans = KMeans(n_clusters=k, random_state=42)

    kmeans.fit(X_scaled)

    sse_scaled.append(kmeans.inertia_)


# Step 3: Plot Elbow Comparison (Scaled vs Unscaled)

plt.figure(figsize=(10, 5))


plt.plot(range(1, 11), sse_unscaled, marker='o', label='Unscaled')

plt.plot(range(1, 11), sse_scaled, marker='s', label='Scaled')

plt.title('Elbow Method (Petal Features Only)')

plt.xlabel('Number of Clusters (k)')

plt.ylabel('SSE (Inertia)')

plt.legend()

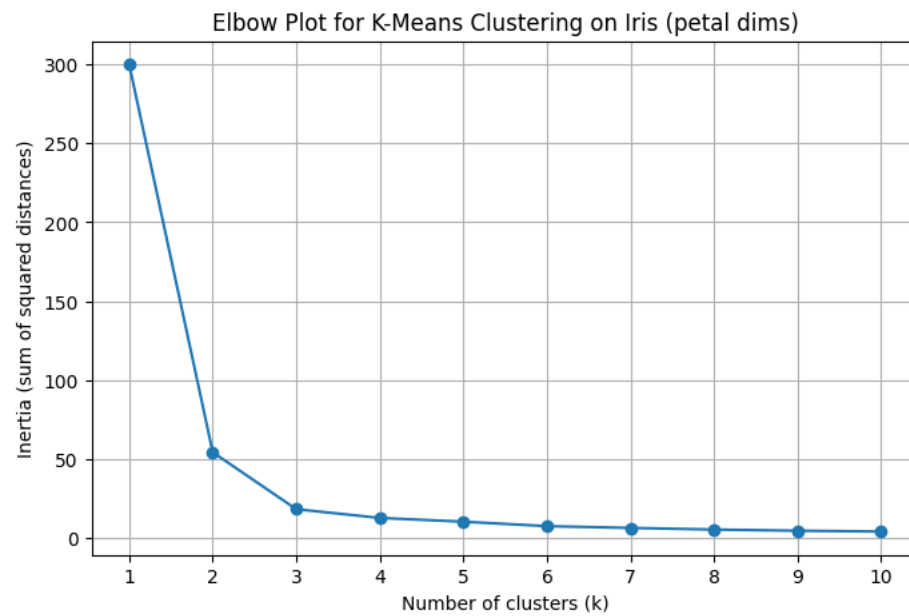
plt.grid(True)

plt.show()

```

Output :

```
k=1, inertia=300.00  
k=2, inertia=54.15  
k=3, inertia=18.05  
k=4, inertia=12.52  
k=5, inertia=10.14  
k=6, inertia=7.31  
k=7, inertia=6.19  
k=8, inertia=5.16  
k=9, inertia=4.41  
k=10, inertia=3.90
```



PROGRAM 11 Implement Dimensionality reduction using Principal Component Analysis (PCA) method.

Screenshot

plates .

Lab 11

Q. Reduce elimination from 2 to 1.

$X_1 = 4 \quad 8 \quad 13 \quad 7$

$X_2 = 11 \quad 4 \quad 15 \quad 14$

Standardize the dataset

$\bar{X}_1 = \frac{4+8+13+7}{4} = 8$

$\bar{X}_2 = \frac{11+4+15+14}{4} = 8.5$

$X'_1 = \begin{bmatrix} 4-8 & 8-8 & 13-8 & 7-8 \\ 11-8.5 & 4-8.5 & 15-8.5 & 14-8.5 \end{bmatrix}$

$X'_1 = \begin{bmatrix} -4 & 0 & 5 & -1 \\ 2.5 & -4.5 & 6.5 & 5.5 \end{bmatrix}$

Covariance matrix

$C = \frac{1}{n-1} X'^T X'$

$C = \frac{1}{3} \begin{bmatrix} 17 & 0 & 15 & -1 \\ 0 & 23 & -33 & 69 \end{bmatrix}$

Let $(C - \lambda I) = 0$

$\begin{vmatrix} 14 - \lambda & -11 \\ -11 & 23 - \lambda \end{vmatrix} = 0$

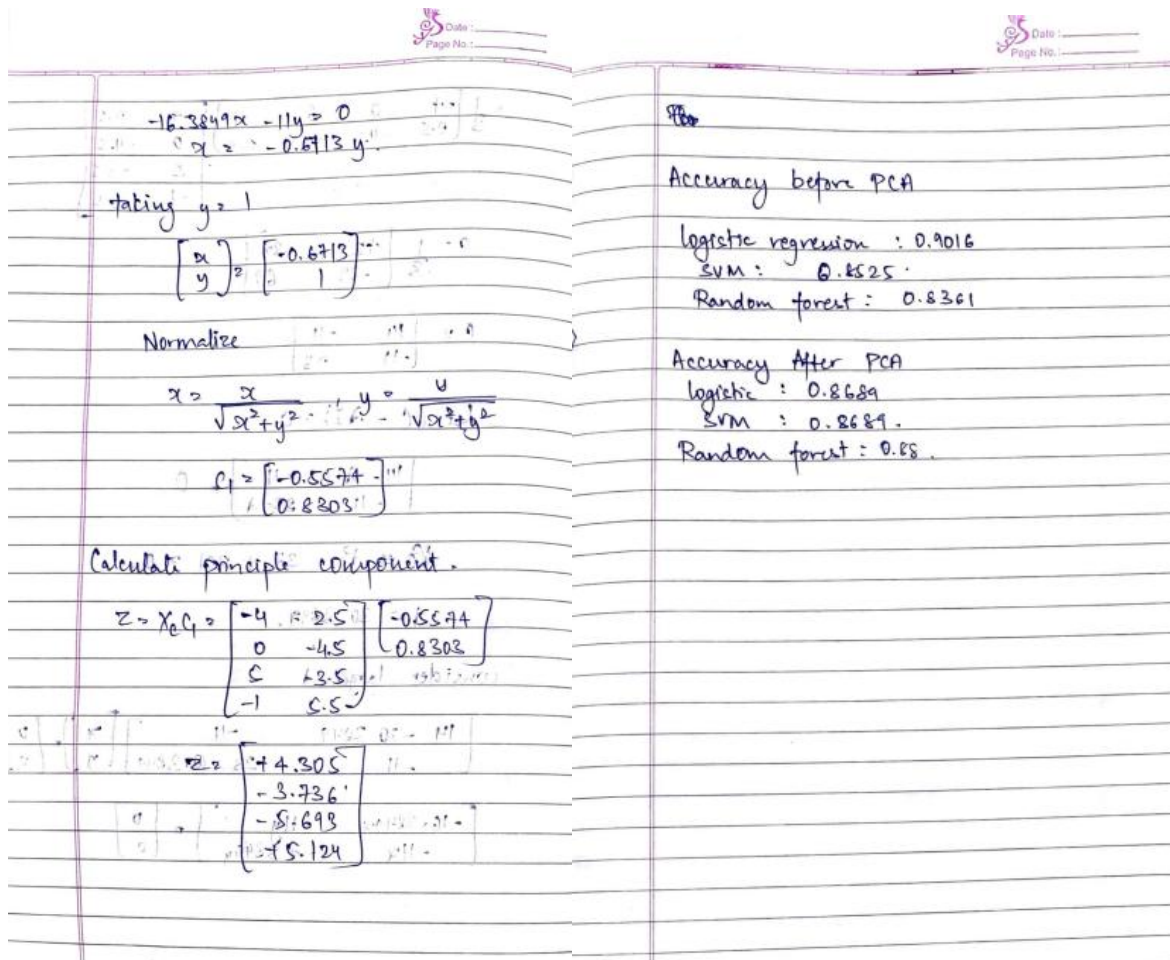
$\lambda^2 - 37\lambda + 201 = 0$

$\lambda_1 = 30.3849, \lambda_2 = 6.6151$

consider larger λ

$\begin{bmatrix} 14 - 30.3849 & -11 \\ -11 & 23 - 30.3849 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$

$\begin{bmatrix} -16.3849x & -11y \\ -11x & -7.3849y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$



Code

import pandas as pd

from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler

from sklearn.compose import ColumnTransformer

from sklearn.pipeline import Pipeline

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.linear_model import LogisticRegression


```

from sklearn.ensemble import RandomForestClassifier

from sklearn.decomposition import PCA

from sklearn.metrics import accuracy_score


# 1. Load data

df = pd.read_csv("heart.csv")


# 2. Label-encode binary text columns

le = LabelEncoder()

for col in ["Sex", "ExerciseAngina"]:

    df[col] = le.fit_transform(df[col])


# 3. Separate features and target

X = df.drop("HeartDisease", axis=1)

y = df["HeartDisease"]


# 4. Build preprocessing pipeline:

# - One-hot for multi-category columns (using sparse_output=False)

# - passthrough the rest

# - then scale everything

cat_cols = ["ChestPainType", "RestingECG", "ST_Slope"]

preprocessor = Pipeline([

    ("onehot", ColumnTransformer([

        ("ohe", OneHotEncoder(sparse_output=False, drop="first"), cat_cols)
    ]))
])

```

```
], remainder="passthrough")),  
("scaler", StandardScaler())  
])
```

5. Apply preprocessing

```
X_proc = preprocessor.fit_transform(X)
```

6. Train/test split

```
X_train, X_test, y_train, y_test = train_test_split(  
    X_proc, y, test_size=0.2, random_state=42  
)
```

7. Define models

```
models = {  
    "SVM": SVC(random_state=42),  
    "LogisticRegression": LogisticRegression(max_iter=1000, random_state=42),  
    "RandomForest": RandomForestClassifier(random_state=42)  
}
```

8. Train & evaluate before PCA

```
print("=== Accuracies BEFORE PCA ===")  
scores_before = {}  
for name, clf in models.items():  
    clf.fit(X_train, y_train)
```

```

preds = clf.predict(X_test)

acc = accuracy_score(y_test, preds)

scores_before[name] = acc

print(f" {name:17s}: {acc:.4f}")

```

9. Apply PCA (retain 95% variance)

```

pca = PCA(n_components=0.95, random_state=42)

X_train_pca = pca.fit_transform(X_train)

X_test_pca = pca.transform(X_test)

print(f"\nPCA retained {pca.n_components_} components, "
      f"explained variance = {pca.explained_variance_ratio_.sum():.4f}\n")

```

10. Train & evaluate after PCA

```

print("=== Accuracies AFTER PCA ===")

scores_after = { }

for name, clf in models.items():

    clf.fit(X_train_pca, y_train)

    preds = clf.predict(X_test_pca)

    acc = accuracy_score(y_test, preds)

    scores_after[name] = acc

print(f" {name:17s}: {acc:.4f}")

```

Output :

```

=== Accuracies BEFORE PCA ===
SVM          : 0.8750
LogisticRegression: 0.8533
RandomForest : 0.8641

PCA retained 13 components, explained variance = 0.9719

=== Accuracies AFTER PCA ===
SVM          : 0.8750
LogisticRegression: 0.8533
RandomForest : 0.8641

```